

Name: Soham Sonar

CWID: A20541266

Table of Contents:

1. Overview
2. Objectives
3. System Architecture
4. Key Components and Responsibilities
5. Network Design
6. Data Flow and Operations
7. Design Decisions and Tradeoffs
8. Future Improvements
9. Conclusion

1. Overview

This design document describes the enhancements to the decentralized Peer-to-Peer Publisher-Subscriber System. Building on the PA3 design, the system now supports **replicated topics** for fault tolerance and performance optimization, and **dynamic topology configuration** to accommodate runtime changes in network nodes.

2. Objectives

- **Performance Optimization:** Reduce latency by replicating topics closer to clients.
 - **Fault Tolerance:** Ensure continuous system operation despite node failures.
 - **Dynamic Adaptability:** Allow seamless addition and removal of nodes in the hypercube topology.
 - **Data Consistency:** Maintain consistency of replicated topics across nodes.
-

3. System Architecture

The enhanced system retains the **hypercube topology** and **Distributed Hash Table (DHT)** but incorporates new features to address performance and fault tolerance:

1) Replicated Topics for performance optimization & fault Tolerance.

2) Dynamic Topology configuration (add and removal of hypernode)

node[000] : Neighbours: [001, 010, 100]

(Data)

Replica

steps: 1) First implement replicas of data in a node to its neighbours.

Create topic ~~from~~ in 000

↓

Def Create topic:

locally saved

send request (replica function - connect(001), (010), (100))

↳ [001, 010, 100]

↳ enter in 001 replica function

Def Replica:

save topic locally in self-replica

↳ enter 010

...

↳ enter 100

operation

complete.

- **Replication Layer:** Extends the DHT with replication capabilities, enabling topics to exist on multiple nodes.
 - **Fault Detector:** Detects node failures and reroutes requests to nodes with replicated topics.
-

4. Key Features

4.1 Replicated Topics

1. Performance Optimization:

- Topics are replicated on nodes physically closer to clients for reduced access latency.
- Replica placement decisions consider factors like network latency, client access frequency, and node load.

2. Fault Tolerance:

- Replicas ensure continuity in topic availability during node failures.
- A synchronization mechanism ensures consistency across all replicas.

3. Consistency Model:

- Adopt a **Primary-Secondary Replication Model**:
 - **Primary Node:** Handles all write operations.
 - **Secondary Nodes:** Handle read operations and synchronize periodically.

4.2 Dynamic Topology Configuration

1. Node Addition:

- New nodes register with the network, receive an ID, and synchronize any topics that hash to their space.
- Topics stored temporarily on other nodes migrate back to the appropriate new node.

2. Node Removal:

- Nodes notify their neighbors of their departure, allowing data to migrate to remaining nodes.
 - Failed nodes are detected by the fault detector, which triggers replica reassignment.
-

5. Network Design

The **hypercube topology** is preserved but enhanced with:

1. **Dynamic Routing:** Routing tables adapt to reflect real-time node availability.
 2. **Redundancy:** Each node maintains references to replicas stored within its neighborhood.
-

6. Data Flow and Operations

6.1 Topic Creation and Assignment

- Topics are hashed to a primary node and assigned to secondary nodes for replication.
- When the primary node fails, a secondary node is promoted to primary.

6.2 Message Routing

- Reads are directed to the nearest replica.
- Writes always go to the primary, with changes propagated asynchronously to replicas.

6.3 Fault Handling

- Failed nodes are detected through periodic heartbeat checks.
- Neighbors of failed nodes update routing tables and reassign topics.

6.4 Dynamic Changes

- Addition and removal of nodes trigger topology updates.
 - Migrating topics maintain seamless access for clients.
-

7. Design Decisions and Tradeoffs

1. Replication Strategy:

- **Decision:** Adopt primary-secondary for consistency.
- **Tradeoff:** Consistency overhead may reduce write performance.

2. Topology Adaptation:

- **Decision:** Dynamic updates based on runtime conditions.
- **Tradeoff:** Increased complexity in routing and fault handling.

3. Fault Tolerance:

- **Decision:** Prioritize availability over strict consistency.
 - **Tradeoff:** Temporary inconsistencies during synchronization.
-

8. Possible Improvements

1. Advanced Replica Placement:

- Use machine learning to predict optimal replica locations based on usage patterns.

2. Enhanced Fault Detection:

- Leverage distributed consensus protocols like Paxos or Raft for fault handling.

3. Dynamic Load Balancing:

- Monitor real-time node usage to redistribute topics dynamically.

4. Elastic Scaling:

- Integrate cloud resources to handle sudden load spikes.
-

9. Conclusion

The enhanced system builds on the foundational hypercube topology and DHT, offering **replicated topics** and **dynamic topology configuration** to meet modern requirements for

scalability, performance, and fault tolerance. Future extensions can focus on predictive optimizations and integration with broader distributed systems.